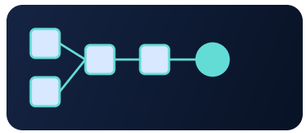
1 Input

Loaded model



VLM · LLM · ViT

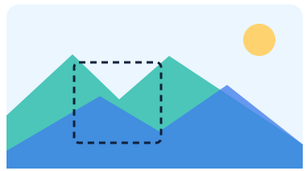


Image and/or text

Question / prompt

Optional pair

Base / clean

Reference input

Native model inputs

2 Configure

Prober + adapter

```
Prober(model)
probe.describe()
```

Model components

Layers · heads · Q/K/V

Token groups

Visual · text · prediction

Target score

Token / class / margin

```
.run(inputs, ...)
```

Optional fit / load

Calibrated lenses

SAE dictionaries

Train + held-out inputs

Save / load an artifact

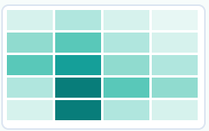
Check capabilities first

3 Tools

Lens

Readout · optional calibrated decoder

Logit · Tuned · Jacobian
Attention Lens · EmbedLens
Patchscope



Layer–token map

Attention

Observation · target-score attribution

Rollout · Grad-CAM · ATTATTR
TAM · Beyond Intuition
DTD/LRP



Attribution map

Causal

Intervention · attribution estimates

Activation / path patching
Knockout · Steering
Attribution / EAP-IG estimates

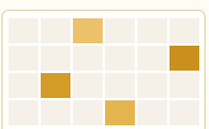


Score changes

SAE

Learned dictionary · feature analysis

ReLU · TopK · Gated · JumpReLU
BatchTopK · Matryoshka · Pursuit
Encode · Reconstruct · Edit



Sparse features

4 Output

Structured outputs

Readouts

Layer / head token scores

Attribution

Maps and effect estimates

Interventions

Changes in answer scores
Changes in generated text

Sparse features

Feature activity and maps
Reconstruction and sparsity

Experiment archive

- Input examples
- Model and method settings
- Analysis results
- Figures and plots
- Example notebooks
- Lenses and SAE dictionaries